

Dhruv Jain

B.Tech CSE (AI & Data Science) | BML Munjal University | Class of 2027
+91 9911850506 | jaindhruv1923@gmail.com | LinkedIn | GitHub | Portfolio

Education

B.Tech, Computer Science Engineering (AI & Data Science)

BML Munjal University, Gurugram

2023 – 2027

Final Year

Skills

Languages & Databases: Python, SQL (PostgreSQL, MySQL), DuckDB, R (Basics), C/C++

Machine Learning & Math: Scikit-learn, EDA, CLV & Churn Modeling, RFM, SHAP, Linear Algebra, Probability & Statistics

BI & Analytics: Power BI (DAX, Power Query, Star Schema), Advanced Excel (Pivot Tables, Solver), Pandas, NumPy

AI & Frameworks: RAG Pipelines, Qdrant Vector DB, Sentence Transformers, Gemini API, FastAPI, Streamlit, REST APIs

Experience

Udaghosh Social Welfare Society (*Data Analyst Intern – Remote*)

July 2026 – October 2026

- Wrote optimized SQL queries (joins, CTEs, window functions, aggregations) to extract, clean, and structure multi-source data into unified, quality-audited reporting datasets.
- Designed interactive Power BI dashboards and automated Excel reporting models (Pivot Tables, dynamic formulas) to track operational KPIs and cut manual reporting turnaround time.
- Performed statistical analysis and monitored KPIs, documenting reports for stakeholder decision-making.

Contentora Media (*MIS Analyst Intern – Remote*)

July 2026 – October 2026

- Maintained structured MIS trackers in Excel to consolidate business research and operational metrics, delivering weekly executive reporting packages.
- Ran systematic data validation, cleaning, and market categorization across research streams to ensure audit-ready data consistency.

Projects

Profitara – Retail BI & Predictive Analytics Pipeline

GitHub → | Live Demo →

Python, SQL, DuckDB, PostgreSQL, Scikit-learn, Power BI, Streamlit, DAX

- Architected an end-to-end retail BI pipeline (ETL to dashboard) over 10,000 transactions across 9 categories (Rs. 66.95L revenue) and 541k records using Python, SQL, DuckDB, and PostgreSQL.
- Built a Random Forest CLV regressor ($R^2 = 0.930$) and a Logistic Regression churn classifier (AUC = 0.91), isolating 361 at-risk accounts for retention.
- Delivered a 13-page Streamlit dashboard and a Power BI star-schema report with K-Means RFM segmentation (silhouette = 0.611, 119 “Champion” customers), 52 Apriori rules, and Isolation Forest discount-abuse detection.

Naukri Saaf – Ghost Job Listing Detection & Analytics

GitHub →

Python, FastAPI, Snorkel, Scikit-learn, TreeSHAP, MySQL 8.0, Power BI, Chrome Extension

- Collected and cleaned 2,851 live job postings from LinkedIn, Indeed, and Glassdoor using Apify and Pandera, identifying 29.3% as ghost postings.
- Applied Snorkel weak supervision (Cohen’s $\kappa = 0.589$) with 5-fold GroupKFold cross-validation. Benchmarked 5 classifiers, with Gradient Boosting reaching 0.72 AUC (0.920 calibrated ROC-AUC) and predictions explained via TreeSHAP.
- Shipped a sub-15 ms FastAPI scoring microservice, an 8-page Power BI dashboard (42-query MySQL workbench), and a real-time Chrome extension that flags ghost listings on job pages.

KAVACH – Security-Governed Agentic AI & Data Platform

GitHub →

Python, FastAPI, Qdrant Vector DB, Sentence Transformers, Gemini API, AST Analysis, pytest, CI/CD

- Built a Python/FastAPI retrieval pipeline with Qdrant semantic search, Sentence Transformers, and the Gemini API for codebase reasoning.
- Authored a deterministic PII and financial-data scanner (Aadhaar/PAN, phones, emails) with 1.00 precision, recall, and F1 to block unsafe content.
- Developed a dependency-graph change-impact analyzer (0.72 avg F1) that maps downstream blast radius, validated by 160+ pytest tests and CI gates.

Certifications

- Microsoft Power BI Data Analyst Track (PL-300 Prep)** – Microsoft / Coursera: Star-Schema Modeling, DAX, Power Query, ETL & Dashboards